

Falco

Cloud native security a observability monitoring

Dive in! →



Petr Michalec

Speaker

Pracuje jako SRE / DevOps v F5 Czech Republic s.r.o.
(Volterra.io, Mirantis, IBM, ...)

n(vi)m lover • maker • golfer • quad fpv pilot

- http://apealive.net/about_petr/
- <http://twitter.com/epcim>



Security monitoring

Účel a požadavky

- ...
- Compliance with security standards:
 - PCI DSS, CIS, GDPR
 - HIPPA, NIST, FIPS, FedRAMP
 - evidence collection
 - data availability and traceability
 - measurement and support resources
 - tools, policies, processes, reporting

Security monitoring

SecOps coverage

- prevent intrusion
- detect intrusion
- configuration enforcement
- compliance checks & reporting
- ...



Traditional approach ~2015

Topics and tools

- network IDS
 - filesystem integrity
 - system/service/user audit/logs
 - data access & encryption
 - security threads, mitigation
 - CVEs, vulnerabilities
 - malicious activity detection
 - auditd, aide
 - PAM, SELinux, AppArmour
 - OSSEC, OpenSCAP, Inspec, ...
 - Enterprise SIEM tools
- SaaS approaching (Snyk, Whitesource, Graylog, Thread stack, ...)

New challenges

Microservices

Containers

- namespace isolation
 - container images
 - 3rd party libraries

Distributed applications

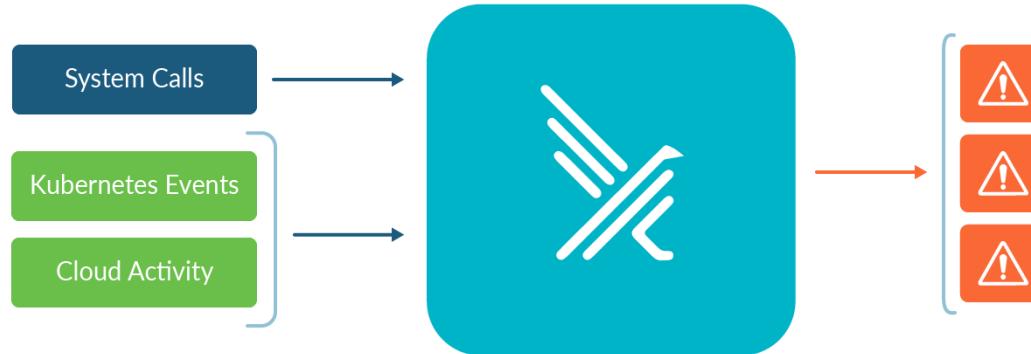
- cloud environments ^n
- 3rd party base images for OS
- 3rd party accessing servers

Distributed data



Falco

May 16, 2016 - Sysdig introducing open source, behavioral security



Falco.org runtime security project detecting unexpected behavior, intrusions, and data theft in real time!

Overview

Falco

- Kernel integration
- Highly granular rules to check for activities involving
 - file and network activity
 - process execution
 - IPC, ...
- Real-time metrics & notification when these rules are violated
- Less complex & faster

There are a million ways a burglar can break into your home, but once they do they're going to steal your jewelry.

...

You only need to detect the things that an attacker does once they have access to a system, rather than all the ways an attacker can gain access to a system.

Comparison to existing approaches

Falco

File integrity monitoring: (checksums)

Watch for any OS activity that is writing to a file of interest, and be alerted in real-time.

Network monitoring (signatures)

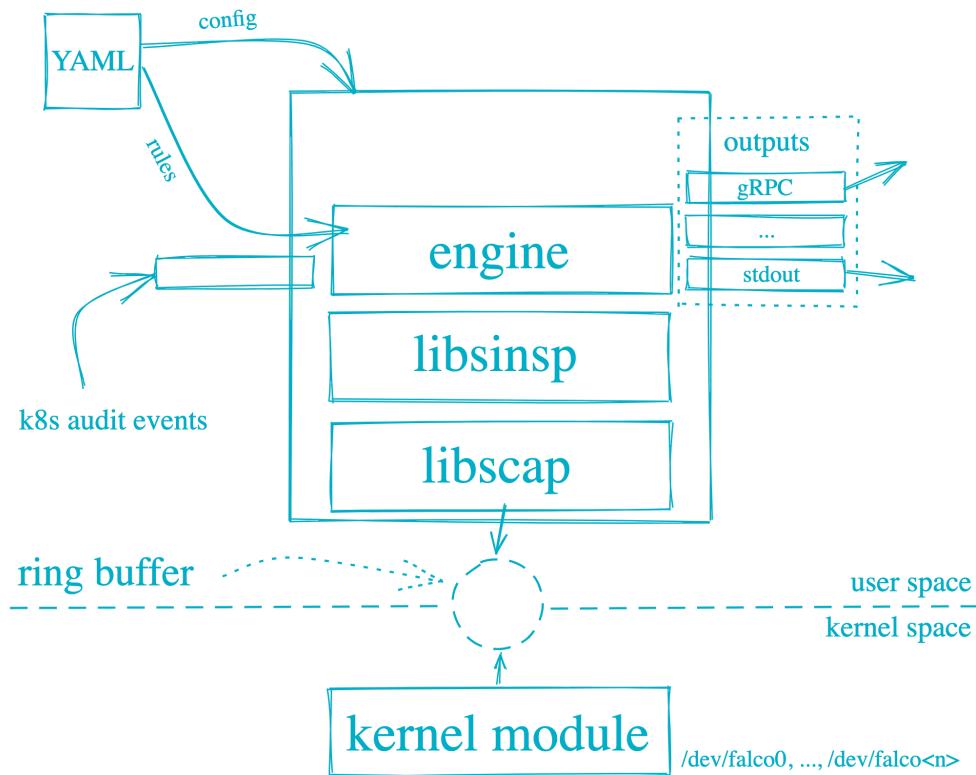
Falco see I/O "from the inside" with an immediate correlation between applications and traffic.

Linux has multiple security modules ~ advanced access control systems with sophisticated policies and concepts. As a result, understanding and configuring them is a rather complex undertaking.

Falco is far simpler to understand and configure, "detection-only".

How it works

Falco architecture



Kernel driver

Falco

- eBPF
- Built-in
- Module `falco.ko` (w/ DKMS)
- Userspace instrumentation (based on PTRACE2)

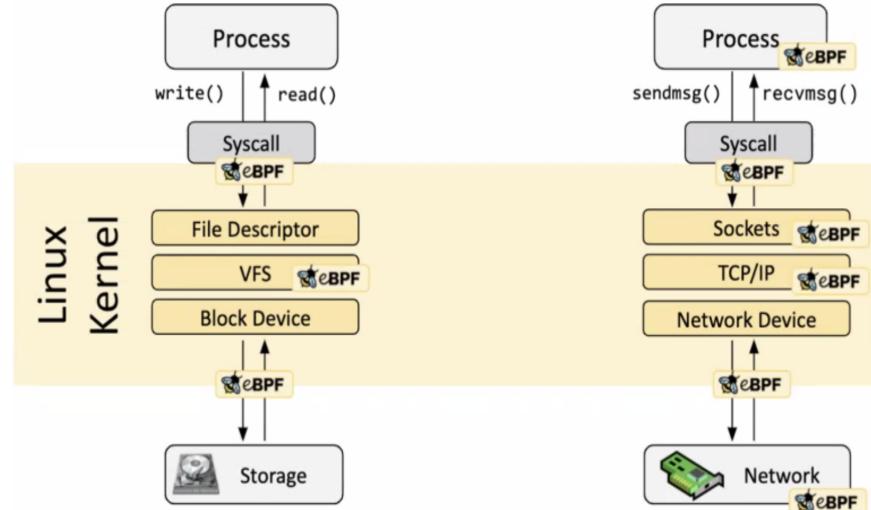
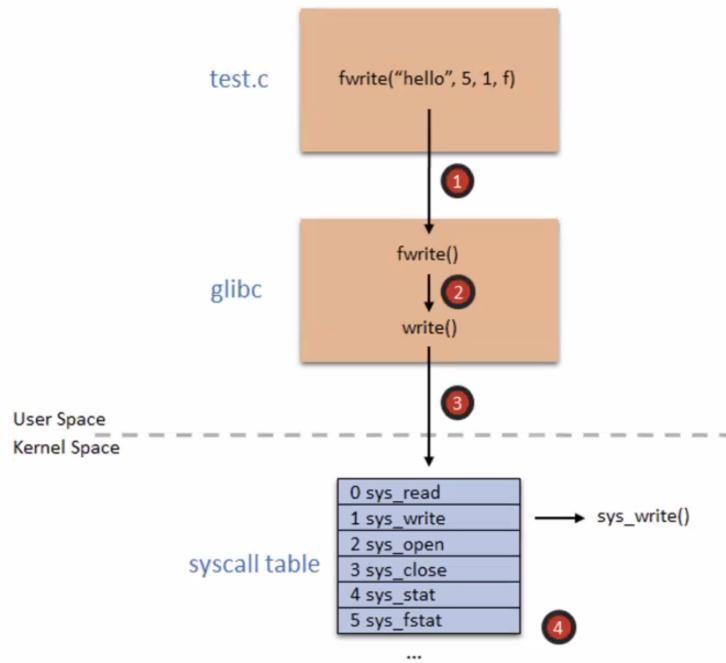
eBPF

- Legacy "Berkeley Packet Filter" (BPF) - technology that among other things allows programs to analyze network traffic (and eBPF is an extended BPF JIT virtual machine in the Linux kernel).
 - raw interface to data link layers
 - permitting raw link-layer packets to be sent and received
 - can run sandboxed programs in a privileged context

BPF is a highly flexible and efficient virtual machine-like construct in the Linux kernel allowing to execute bytecode at various hook points in a safe manner. It is used in number of Linux kernel subsystems (networking, tracing, security (sandboxing))"

SysCalls

eBPF



Why?

eBPF

Enhanced Telemetry Collection -> annotation

- kernel and syscall attributes
- socket info

Performance

- lower resource impact (net, file, proc)
- avoid transfer of all audit data to userspace
- real time processing

eBPF Verifier verifies the safety of eBPF programs

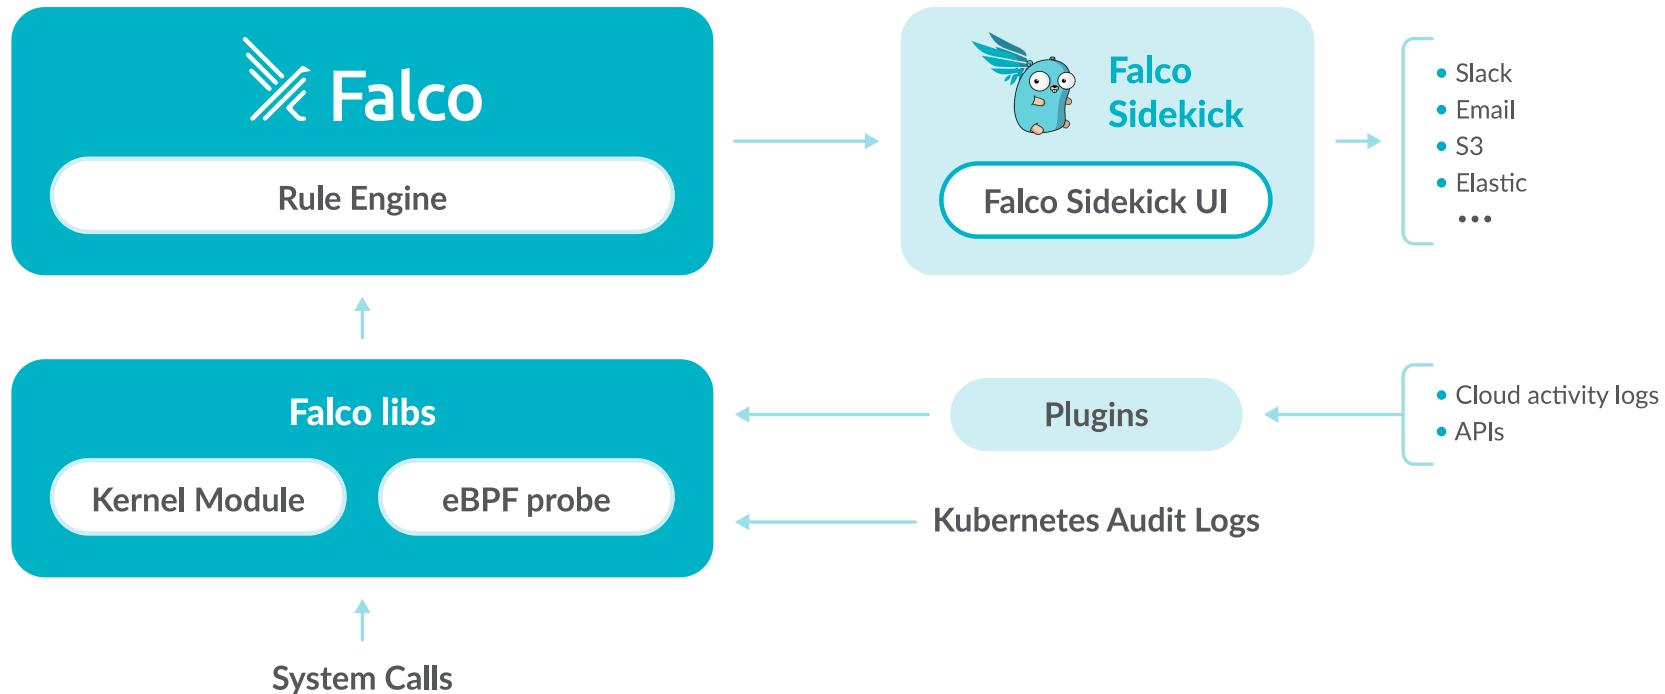
Field Class: process

Additional information about the process and thread executing the syscall event.

Name	Type	Description
proc.pid	INT64	the id of the process generating the event.
proc.exe	CHARBUF	the first command line argument (usually the executable name or a custom one).
proc.name	CHARBUF	the name (excluding the path) of the executable generating the event.
proc.args	CHARBUF	the arguments passed on the command line when starting the process generating the event.
proc.env	CHARBUF	the environment variables of the process generating the event.
proc.cmdline	CHARBUF	full process command line, i.e. proc.name + proc.args.
proc.exeline	CHARBUF	full process command line, with exe as first argument, i.e. proc.exe + proc.args.
proc.cwd	CHARBUF	the current working directory of the event.
proc.threads	UINT32	the number of threads that the process generating the event currently has, including the main process thread.
proc.childs	UINT32	the number of child threads that the process generating the event currently has. This excludes the main process thread.
proc.ppid	INT64	the pid of the parent of the process generating the event.
proc.pname	CHARBUF	the name (excluding the path) of the parent of the process generating the event.

Deployment

Falco components



Deployment

Configuration

Deployment

- Falco-sidekick, prom. exporter
- Falco-sidekick UI
- Grafana dashboards
- ..., SysFlow, ELK
- ..., Plugins

Daemonset

- Falco

*(only `falco-driver-loader` needs to be run with
`securityContext: privileged`)*

What to enable?

- driver-loader (DKMS, private builds)
- docker, containerd, cri-o
- w/k8s metadata, audit
- custom rules
- readiness, maxBurst, eventDrops
- priority/severity level

Language

Syntax, github.com/falcosecurity/charts/falco/rules

Macros

- name (identifier)
- condition (filter)

Lists

- name (identifier)
- items:

Rules

- `name` (identifier)
- desc
- `condition` (filter expression, macro)
- `output` (formated message with **core details**)
- priority (severity of rule)
- exceptions (new)

Primitives

Rules

Shell executed in container

```
container.id != host and proc.name = bash
```

Overwirite system bins

```
fd.directory in (/bin, /bin/sbin, /usr/bin, /usr/sbin)  
and write
```

Container namespace change

```
evt.type = setns and not proc.name in (docker)
```

Process access camera

```
evt.type = open and fd.name = /dev/video0 and not proc.name in (skype, zoom, webex)
```

Macros & Lists

Rules

```
- list: _container_engine_binaries
  items: [dockerd, containerd, containerd-shim, "runc:[0:PARENT]", "runc:[1:CHILD]", "runc:[2:INIT]"]

- macro: docker_authorized_binaries
  condition: >
    proc.name in (_container_engine_binaries)
    or proc.pname in (_container_engine_binaries)

  " [CVE-2019-11246 on Mitre](https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2019-11246)"

- macro: safe_kubectl_version
  condition: (
    jevt.value[/useragent] startswith "kubectl/v1.20" or
    jevt.value[/useragent] startswith "kubectl/v1.19"
    ...
  )
```

Filesystem integrity

Rules

```
- rule: Detect Write Below /etc/hosts
  desc: an attempt to write to /etc/hosts file (CVE-2020-8557)
  condition: open_write and container and fd.name=/etc/hosts
  output: "File /etc/hosts opened for writing (user=%user.name command=%proc.cmdline parent=%proc.pname \
            pc cmdline=%proc.pc cmdline file=%fd.name program=%proc.name gparent=%proc.aname[2] \
            ggp parent=%proc.aname[3] ggg parent=%proc.aname[4] container_id=%container.id image=%container.image.repository"
  priority: ERROR
  tags: [filesystem, mitre_persistence]
```

Detect su, sudo

Rules

```
- rule: Detect su or sudo
  desc: detect sudo activities
  condition: >
    spawned_process
    and activity_under_escalated_privilege
    and not in (sre_known_infraops_actions, sre_known_managed_cloud_actions)
  output: >
    Privilege escalation activity (user=%user.name auser=%user.loginname command=%proc.cmdline ppid=%proc.ppid apid=%proc.apid priority: WARNING
  tags: [process, sudo, su]

- macro: activity_under_escalated_privilege
  condition: >
    proc.name in (sudo, su)
    or proc.pname in (sudo, su)
    or proc.aname[1] in (sudo, su)
    ...
  ...

- rule: Privilege escalation
  condition: sf.pproc.uid != 0 and sf.proc.uid = 0 and not entrypoint
```

Compromised server process

Rules

HitchSQL injection attack?

```
condition: spawn_process and proc.name = mysql and not proc_is_new
```

```
- macro: spawn_process
  condition: syscall.type = execve

- macro: proc_is_new
  condition: proc.duration <= 5000000000
```

Howto rules

What we tweaked?

```
# override upstream defined macro
- macro: user_known_read_sensitive_files_activities
  condition: >
    (sre_authorized_activities)
```

```
rules/
├── falco_rules.preload.yaml
├── falco_rules-10-exceptions.yaml
├── falco_rules-20-security.yaml
├── falco_rules-30-apps.yaml
├── falco_rules-40-fim.yaml
└── falco_rules-50-cve.yaml
```

- macro: `failed_k8s_annotation`
- macro: `sre|host|infra_authz_activities`
- macro: `sre|aws|gcp_known_vendoractions`
- macro: `sre_known_ports`
- list: `sre|aws|gcp_known_commands`
- list: ...

Falco sidekick

Integrations

<https://github.com/falcosecurity/falcosidekick>

```
alertmanager:  
  hostport: http://{{domain or ip}}:{{port}}  
  minimumpriority: "error"          # emergency|alert|critical|error|warning|notice|informational|debug  
  endpoint: "/api/v2/alerts"  
  expiresafter: "900"
```

Slack • Rocketchat • Mattermost • Teams • Datadog • Discord • AlertManager • Elasticsearch • Loki • NATS • STAN (NATS Streaming) • Influxdb • AWS Lambda • AWS SQS • AWS SNS • AWS CloudWatch • AWS S3 • SMTP (email) • Opsgenie • StatsD • DogStatsD • Webhook • Azure Event Hubs • Prometheus • GCP PubSub • GCP Storage • Google Chat • Apache Kafka • PagerDuty • Kubeless • OpenFaaS

Dashboards and alerting

The screenshot displays the Falcosidekick UI dashboard, which includes a header with navigation tabs (Dashboard, Events, Info), a search bar, and a summary of total events across various categories. Below this is a table listing recent alerts and log entries, each with details like timestamp, source, priority, rule, output, and associated tags.

Timestamp	Source	Priority	Rule	Output	Tags
2022/05/11 12:18:565	syscalls	Critical	Polkit Local Privilege Escalation Vulnerability (CVE-2021-4034)	Detect Polkit pkexec Local Privilege Escalation Exploit (CVE-2021-4034) (user=%user.loginname uid=%user.loginuid command=%proc.cmdline args=%proc.args) proc.args proc.cmdline %proc.cmdline user.loginname %user.loginname user.loginuid %user.loginuid	process mitre_privilege_escalation
2022/05/11 12:18:12:561	syscalls	Warning	Mount Launched in Privileged Container	Mount was executed inside a privileged container (user=%user.name user_loginuid=%user.loginuid command=%proc.cmdline %container.info image=%container.image.repository %container.image.tag) container.image.repository %container.image.repository container.image.tag %container.image.tag container.info %container.info proc.cmdline %proc.cmdline user.loginuid %user.loginuid user.name %user.name	container cis mitre_lateral_movement
2022/05/11 12:17:58:555	syscalls	Warning	Create Symlink Over Sensitive Files	Symlinks created over sensitive files (user=%user.name user_loginuid=%user.loginuid command=%proc.cmdline target=%evt.arg.target linkpath=%evt.arg.linkpath parent_process=%proc.pname) evt.arg.linkpath %evt.arg.linkpath evt.arg.target %evt.arg.target proc.cmdline %proc.cmdline proc.pname %proc.pname user.loginuid %user.loginuid user.name %user.name	file mitre_exfiltration
2022/05/11 12:17:40:538	syscalls	Notice	Launch Ingress Remote File Copy Tools in Container	Ingress remote file copy tool launched in container (user=%user.name user_loginuid=%user.loginuid command=%proc.cmdline parent_process=%proc.pname container_id=%container.id container_name=%container.name image=%container.image.repository %container.image.tag) container.id %container.id container.image.repository %container.image.repository container.image.tag %container.image.tag container.name %container.name k8s.ns.name %k8s.ns.name k8s.pod.name %k8s.pod.name proc.cmdline %proc.cmdline proc.pname %proc.pname user.loginuid %user.loginuid user.name %user.name	network process mitre_command_and_control
2022/05/11 12:17:38:536	syscalls	Informational	System user interactive	System user ran an interactive command (user=%user.name user_loginuid=%user.loginuid command=%proc.cmdline container_id=%container.id image=%container.image.repository) container.id %container.id container.image.repository %container.image.repository proc.cmdline %proc.cmdline user.loginuid %user.loginuid user.name %user.name	users mitre_remote_access_tools
2022/05/11 12:17:29:528	syscalls	Notice	DB program spawned process	Database-related program spawned process other than itself (user=%user.name user_loginuid=%user.loginuid program=%proc.cmdline parent=%proc.pname container_id=%container.id image=%container.image.repository) container.id %container.id container.image.repository %container.image.repository proc.cmdline %proc.cmdline proc.pname %proc.pname user.loginuid %user.loginuid user.name %user.name	process database mitre_execution

Elasticsearch

Record detail

⌚ output_fields	<pre>{ "user.uid": 2201, "proc.cmdline": "uname -o", "proc.ppid": 6271, "proc.env": "XDG_SESSION_ID=5471 HOSTNAME=master-0 SELINUX_ROLE_REQUESTED= TERM=xterm-256color SHELL=/bin/bash HISTSIZE=1000 S SH_CLIENT=10.54.80.50 51390 22 SELINUX_USE_CURRENT_RANGE= SSH_TTY=/dev/pts/0 USER=vesop MAIL=/var/spool/mail/vesop PATH=/usr/loc al/bin:/usr/bin:/usr/local/sbin:/usr/sbin PWD=/home/vesop SELINUX_LEVEL_REQUESTED= HISTCONTROL=ignoredups SHLVL=1 HOME=/home/ves op DNSServiceIP=10.3.0.10 LOGNAME=vesop SSH_CONNECTION=10.54.80.50 51390 10.63.51.10 22 KUBERNETES_VERSION=v1.15.0 XDG_RUNTIME_D IR=/run/user/2201 _=/usr/bin/uname", "proc.ename[2]": "bash", "container.id": "host", "proc.ename[3]": "sshd", "proc.pid": 6273, "proc.ename[4]": "sshd", "user.loginuid": 2201, "group.gid": 2201, "proc.pcmdline": "bash", "evt.time": 1654426400432897300 }</pre>
⌚ severity	warning
⌚ site	ty8-tky
⚠ source	syscall
⌚ source_type	kafka
⌚ stream	stdout
⌚ tag	kube.app.falco.falco
⌚ tags	group, process, user

Falco Audit in Grafana

Volterra Security Services / Falco Audit

Last 30 minutes

Audit Events

time	severity	message	Cluster	hostname	namespace	pod	procCmdline	fdName	id
2021-12-13 20:22:41	warning	Detect File Permission or Ownership Change	gc01-cle-int-ves-io	ip-172-16-130-194.us-east-2.compute.internal		-	chmod 700 /tmp/awsagent.8k0fdvcy		KaQ_tX0Bsryk
2021-12-13 20:22:40	warning	Clear Log Activities	gc01-cle-int-ves-io	ip-172-16-130-194.us-east-2.compute.internal		-	bash -c /bin/sleep \${((RANDOM % 3000) + 1)}; rm -f /var/log/awsagent-update.log; umask 037 && /opt/aws/awsagent/bin/update > /var/log/awsagent-update.log 2>&1	/var/log/awsagent-update.log	KKQ_tX0Bsryk
2021-12-13 20:21:50	warning	Clear Log Activities	gc01-cle-int-ves-io	ip-172-16-132-140.us-east-2.compute.internal		-	bash -c /bin/sleep \${((RANDOM % 3000) + 1)}; rm -f /var/log/awsagent-update.log; umask 037 && /opt/aws/awsagent/bin/update > /var/log/awsagent-update.log 2>&1	/var/log/awsagent-update.log	5Bw-tX0BRw1
2021-12-13 20:18:29	error	Write below rpm database	gc01-cle-int-ves-io	ip-172-16-147-217.us-east-2.compute.internal		-	python	/var/lib/rpm/.dbenv.lock	oRw7tX0BRw
2021-12-13 20:18:28	error	Write below rpm database	gc01-cle-int-ves-io	ip-172-16-147-217.us-east-2.compute.internal		-	python	/var/lib/rpm/__db.003	HqQ7tX0Bsryl
2021-12-13 20:18:28	error	Write below rpm database	gc01-cle-int-ves-io	ip-172-16-147-217.us-east-2.compute.internal		-	python	/var/lib/rpm/__db.002	HaQ7tX0Bsryl

Plugins

Added recently (>= v0.31)

External sources

- API boundaries, hardly extensible
- Falco must expose a web server
- TLS to manage
- Doesn't work with managed K8s

- K8s audit
- AWS CloudTrail
- JSON
- coming (okta, github, docker, seccompagent)

Features

- dynamic shared libraries -> any language
- allows falco to collect and extract fields from streams of events
- source / extractor plugins

K8s audit rules

<https://github.com/falcosecurity/plugins/tree/master/plugins/k8saudit>

An attempt to start a pod using the host pid NS.

```
condition: kevt and pod and kcreate  
and ka.req.pod.host_pid intersects (true)
```

Updated role binding

```
condition: kevt  
and clusterrolebinding  
and kcreate and ka.req.binding.role=cluster-admin
```

Detect pod starting a privileged container

```
condition: kevt  
and pod  
and kcreate  
and ka.req.pod.containers.privileged intersects (true)  
and not ka.req.pod.containers.image.repository  
in (falco_privileged_images)
```

Credentials in configmap

```
- macro: contains_private_credentials  
condition: >  
(ka.req.configmap.obj contains "access_key" or  
ka.req.configmap.obj contains "access-key" or  
ka.req.configmap.obj contains "token" or  
ka.req.configmap.obj contains "secret" or  
ka.req.configmap.obj contains "pass")
```

CloudTrail

Plugin

```
- rule: Console Login Without MFA
  desc: Detect a console login without MFA.
  condition:
    ct.name=="ConsoleLogin" and not ct.error exists
    and ct.user.identitytype!="AssumedRole"
    and json.value[/responseElements/ConsoleLogin]=="Success"
    and json.value[/additionalEventData/MFAUsed]=="No"
  output:
    Detected a console login without MFA
    (requesting user=%ct.user,
     requesting IP=%ct.srcip,
     AWS region=%ct.region)
  priority: CRITICAL
```

What next?

```
- rule: Pet detection, custom plugin
  condition: video.entities[animal] > 0
```

Sysflow.io

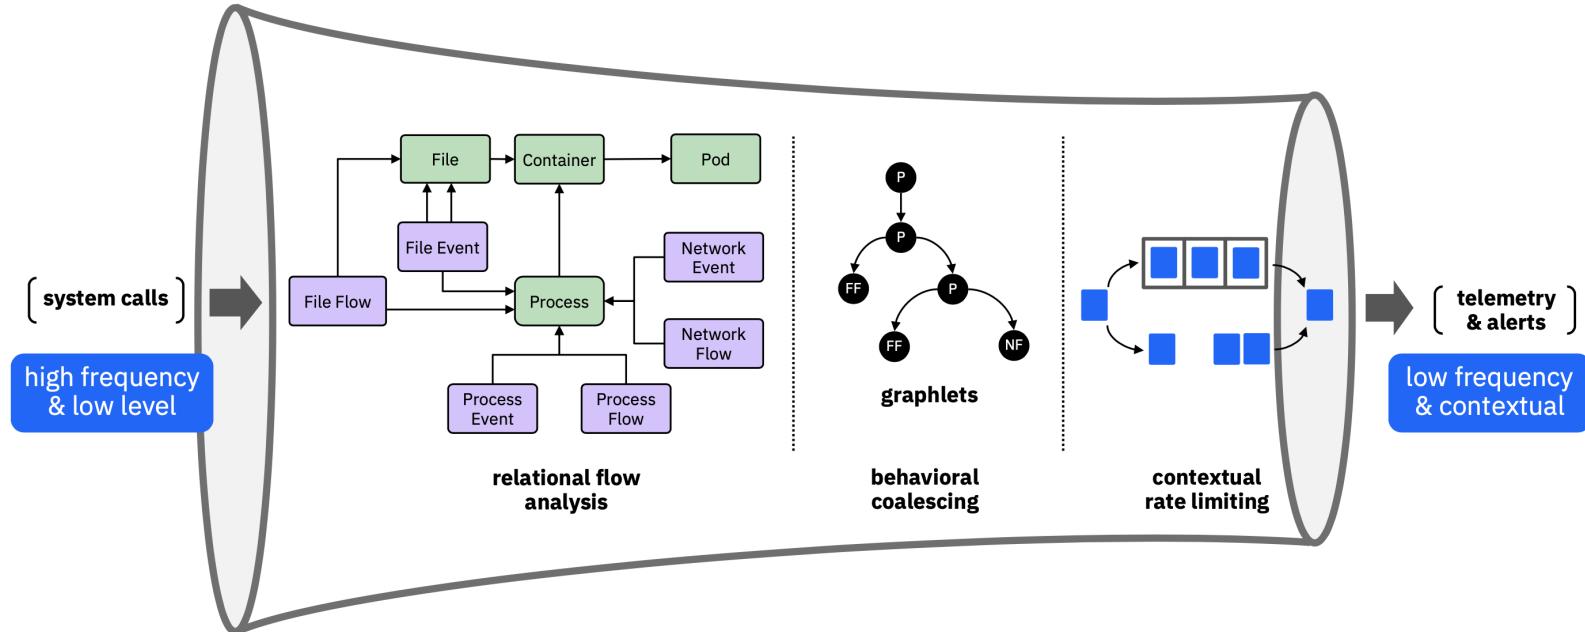
```
- rule: Impair Defenses: Disable or Modify System Firewall
  desc: Detects disabling security tools
  condition: sf.opflags = EXEC and
    ((sf.proc.name in (service_cmds) and
      sf.proc.args pmatch (security_services) and sf.proc.args pmatch (stop_cmds)) or
    (sf.proc.name = setenforce and sf.proc.args = '0'))
  prefILTER: [PE]

- rule: Large network data transfer with database endpoint
  condition: ( sf.opflags contains RECV and sf.net.dport = 3306 and sf.flow.rbytes > 1024 ) or
    ( sf.opflags contains SEND and sf.net.sport = 3306 and sf.flow.wbytes > 1024 )
  prefILTER: [NF]

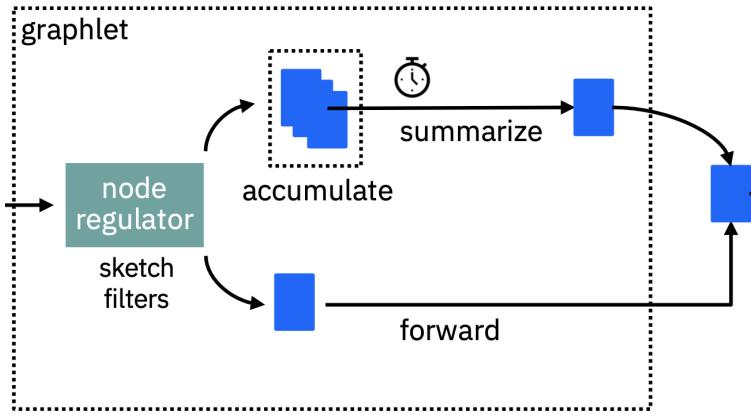
- rule: Privilege escalation
  condition: sf.pproc.uid != 0 and sf.proc.uid = 0 and not entrypoint
```

SysFlow.io

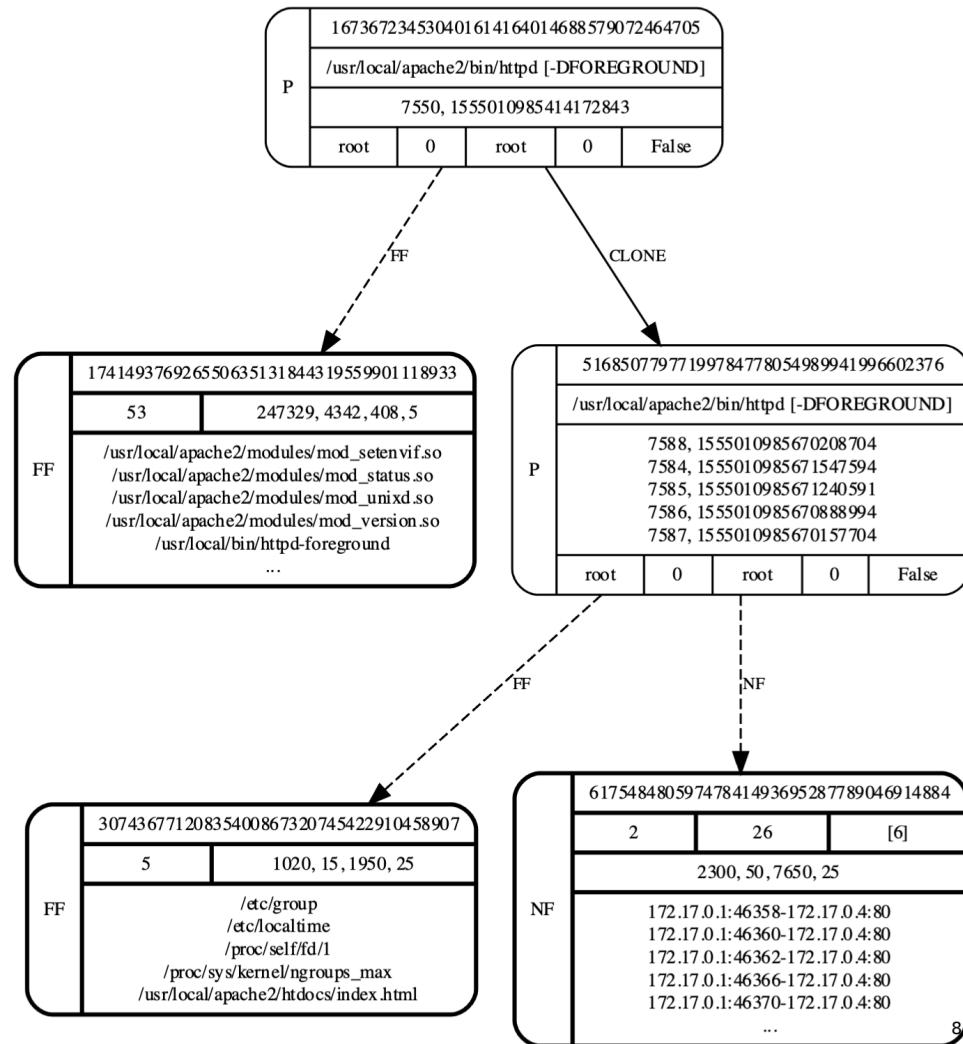
cloud-native system telemetry framework



SysFlow.io



- Rate modulation
- Node-level regulators
 - HyperLogLog sketch
 - Count-min sketch



Learn More

[Documentation](#) · [GitHub](#) · [Blog](#)

- SysFlow is a cloud-native system telemetry framework that enables the creation of security analytics on a scalable, pluggable open-source platform
 - [SysFlow telemetry](#)
 - [SysFlow & Sidekick analytics PoC](#)
 - [SysFlow policies & examples](#)
- [Plugin Pet surveillance with falco PoC](#)
- [Employ AI/ML for anomaly detection](#)

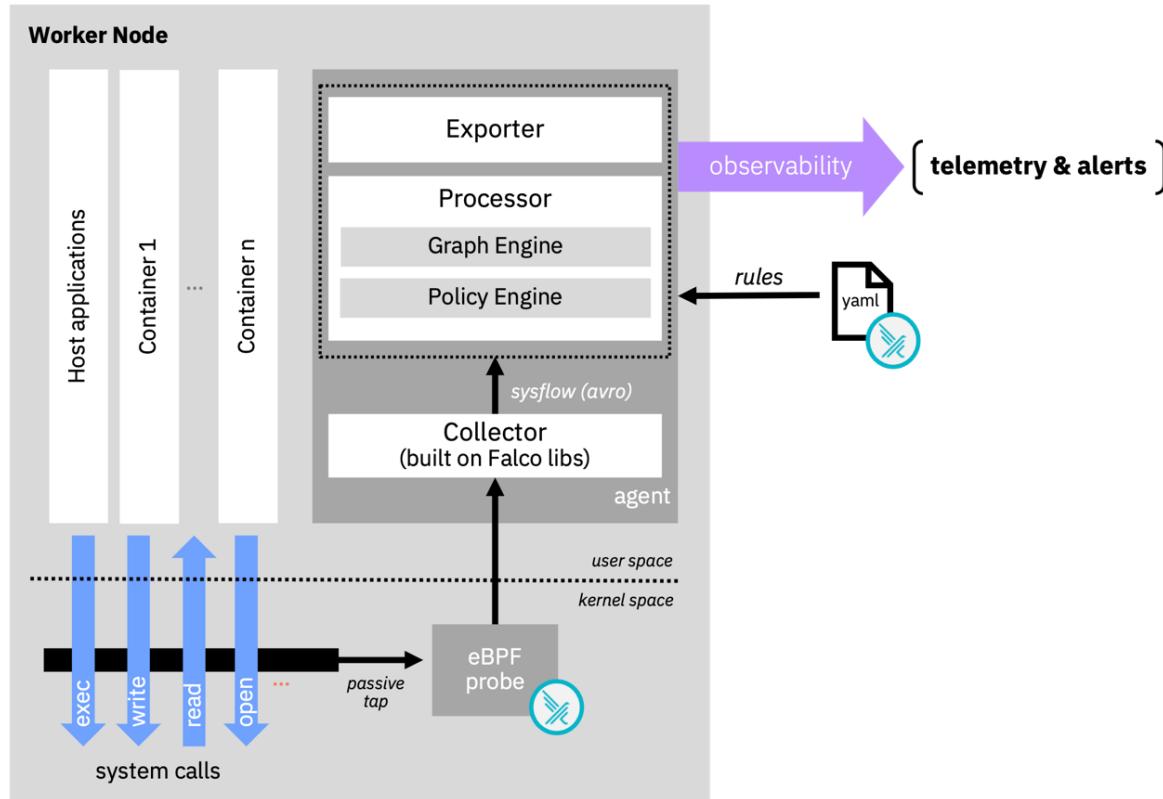
END

Backup slides

Falco architectural overview

Backup slides

SysFlow architectural overview



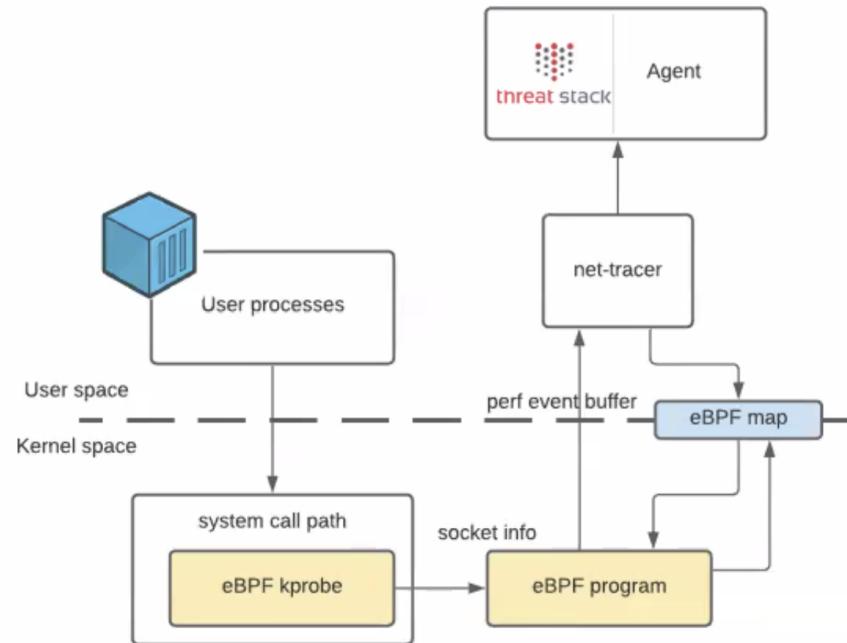
Backup slides

eBPF network observability

The agent attaches eBPF programs to kprobes to trace TCP (connect/accept) and UDP (connect) activity.

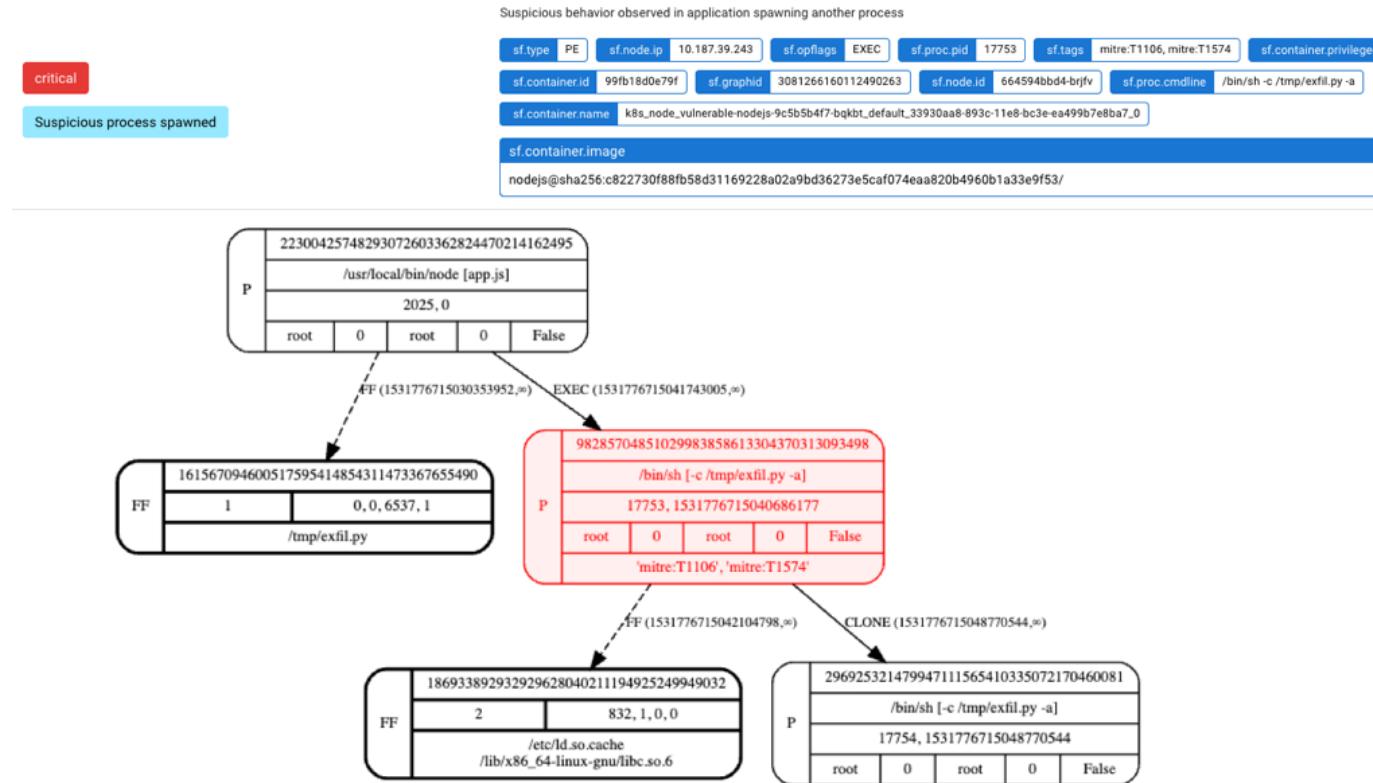
Uses telemetry from eBPF sensor to augment audit events.

Additional DNS sensor provides reverse DNS information



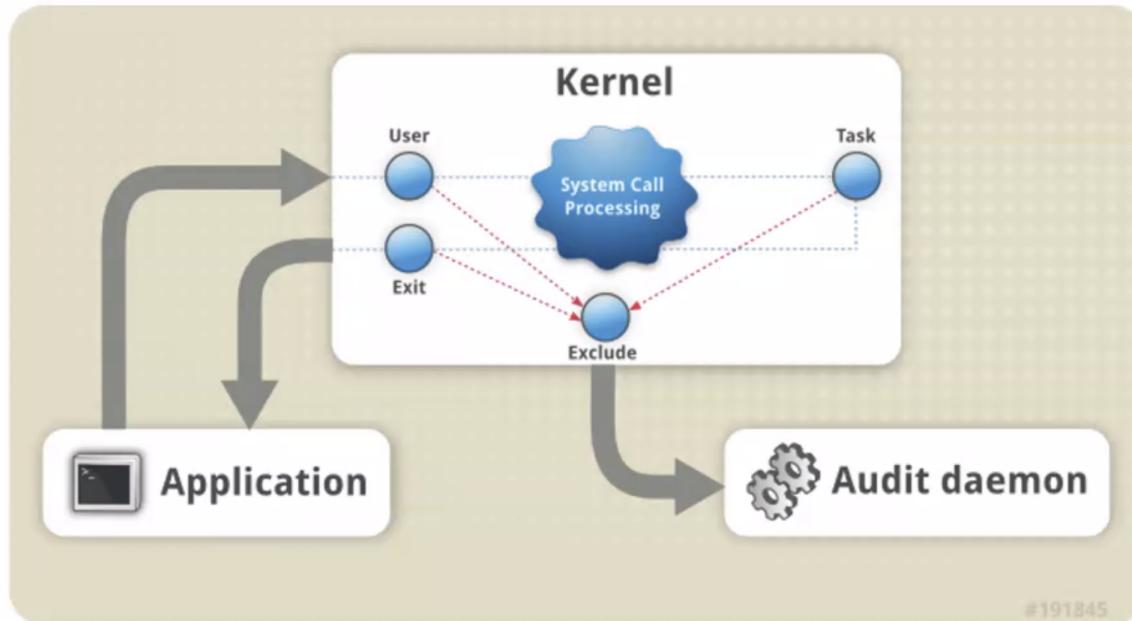
Backup slides

SysFlow.io integration to sidekick



Backup slides

AuditD comparison



Backup slides

Inbound ssh rule

```
- rule: Inbound SSH Connection
  desc: Detect Inbound SSH Connection
  condition: >
    ((evt.type in (accept,listen) and evt.dir=<) or
     (evt.type in (recvfrom,recvmsg))) and ssh_port
     and not is_kubernetes
  output: >
    Inbound SSH connection (user=%user.name client_ip=%fd.cip client_port=%fd.cport server_ip=%fd.sip)
  priority: WARNING
  tags: [ssh, network]
```

Backup slides

K8s audit -> plugin

- Removed K8S audit logs from Falco [#1952] (<https://github.com/falcosecurity/falco/pull/1952>)
- Now under plugins: <https://github.com/falcosecurity/plugins>

```
- rule: Attach/Exec Pod
desc: Detect any attempt to attach/exec to a pod
condition: |
  kevt_started and pod_subresource and kcreate and ka.target.subresource in (exec,attach)
  and not user_known_exec_pod_activities

- list: falco_hostpid_images
items: []

- rule: Create HostPid Pod
desc: Detect an attempt to start a pod using the host pid namespace.
condition: |
  kevt and pod and kcreate and ka.req.pod.host_pid intersects (true)
  and not ka.req.pod.containers.image.repository in (falco_hostpid_images)
```

Backup slides

Rules from helm chart

<https://github.com/falcosecurity/falco/tree/master/rules>

```
rules-traefik.yaml: |-  
  - macro: traefik_consider_syscalls  
    condition: (evt.num < 0)  
  
  - macro: app_traefik  
    condition: container and container.image startswith "traefik"  
  
  # Restricting listening ports to selected set  
  
  - list: traefik_allowed_inbound_ports_tcp  
    items: [443, 80, 8080]
```

Dashboards and alerting

